AI Embedded

How AI-powered hardware is reshaping industries **- in depth case study**

Bhavishya Pandit

AI + Hardware: The Perfect Fusion!

AI is revolutionizing hardware in exciting new ways. In this carousel, you'll learn about:

Intro

- How is AI integerated in Hardware

Applications

- AI-optimized processors
- AI in Defense
- AI in IOT Devices
- AI in Edge Computing
- AI and Robotics

How Is AI Integrated Into Hardware?

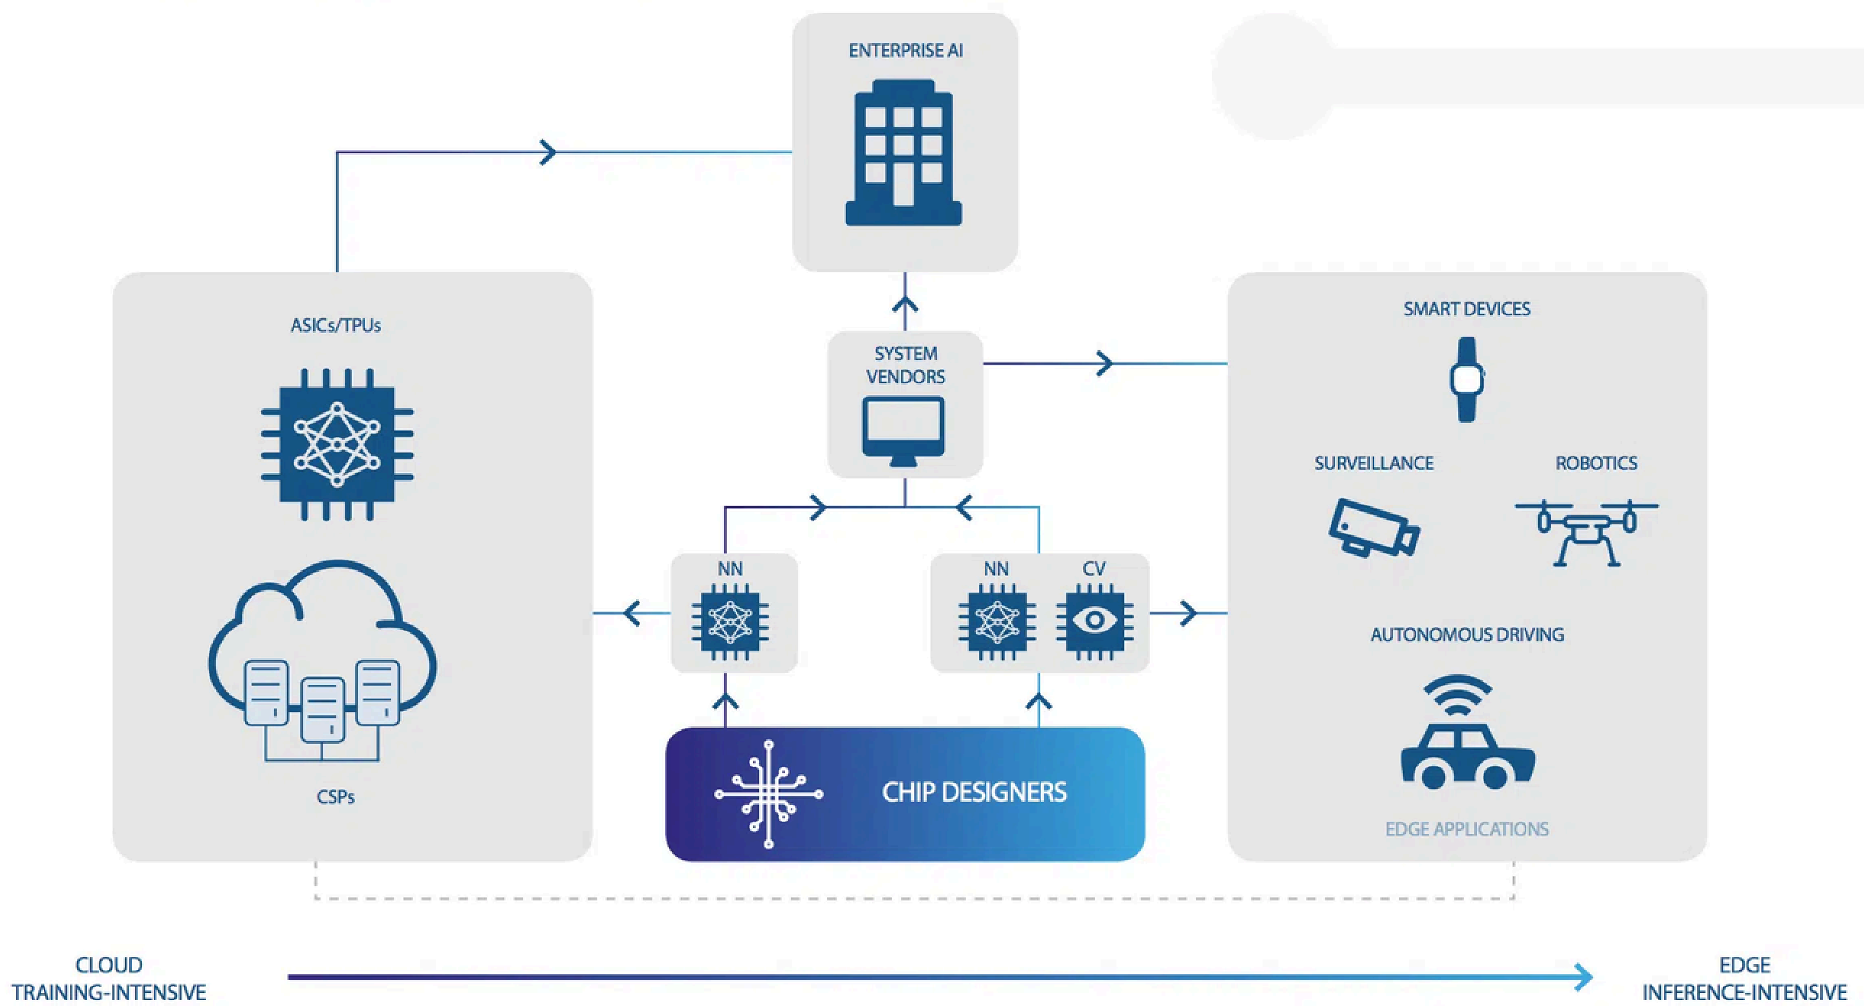

Source

1. **Model development:** Everything starts with training an AI model using large datasets. This is usually done on powerful cloud servers using frameworks like TensorFlow or PyTorch.

2. **Model optimization:** Once trained, the model is too big or slow for everyday hardware. So, it's optimized shrunk down using techniques like quantization. Tools like TensorRT or CoreML make the model smaller, faster, and ready for deployment.

3. **Hardware mapping**

The optimized model is then mapped to specialized hardware:

- CPUs for general tasks
- GPUs for parallel computations
- NPUs/TPUs for dedicated AI workloads
- The right chip is chosen based on what the AI needs to do.

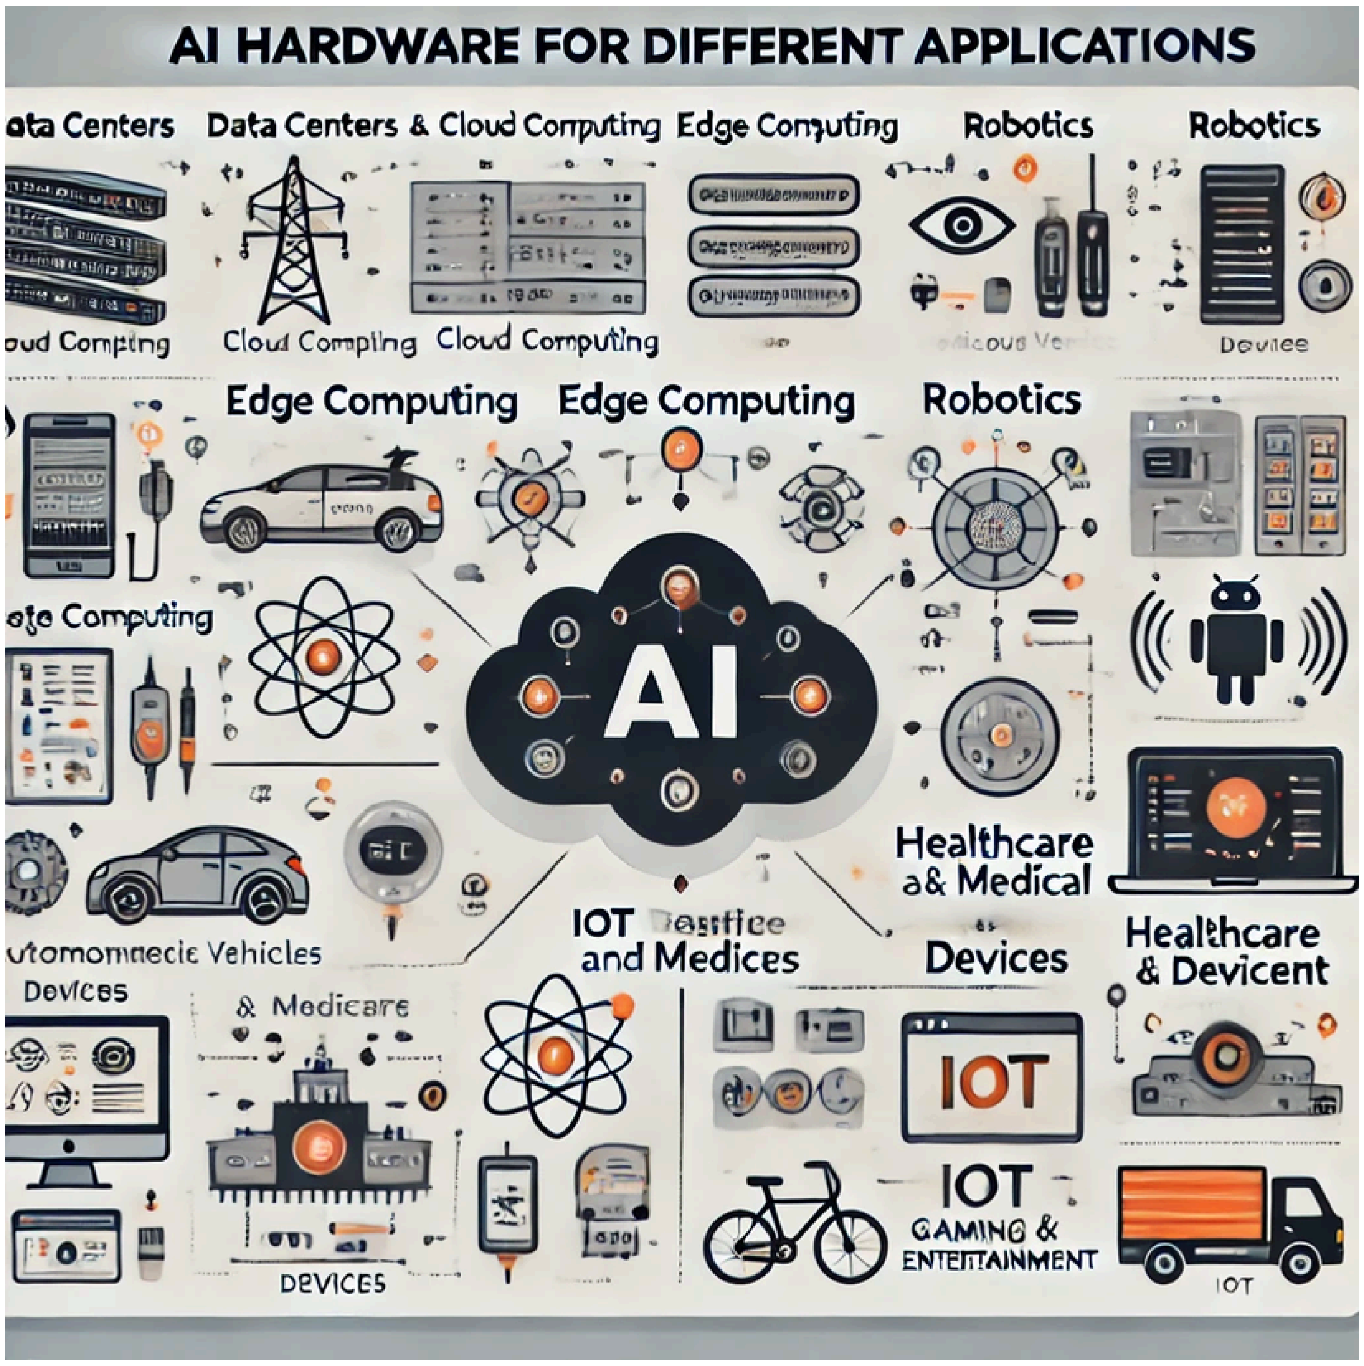

Source

4. **Edge integration**: Now, the hardware with the AI model is embedded into the end device like a phone, camera, drone, or smart speaker. The AI runs directly on the device, which is known as edge computing.

5. **Runtime execution**: When in use, the device collects real-time data (like images, voice, or sensor input), which the AI model processes instantly. No cloud calls this makes everything faster and more private.

6. **Feedback loop**: Some systems include a feedback loop. That means they either learn over time (on-device learning) or send data back to the cloud to improve the model, which can then be updated and reloaded.

1. AI Optimized Processors

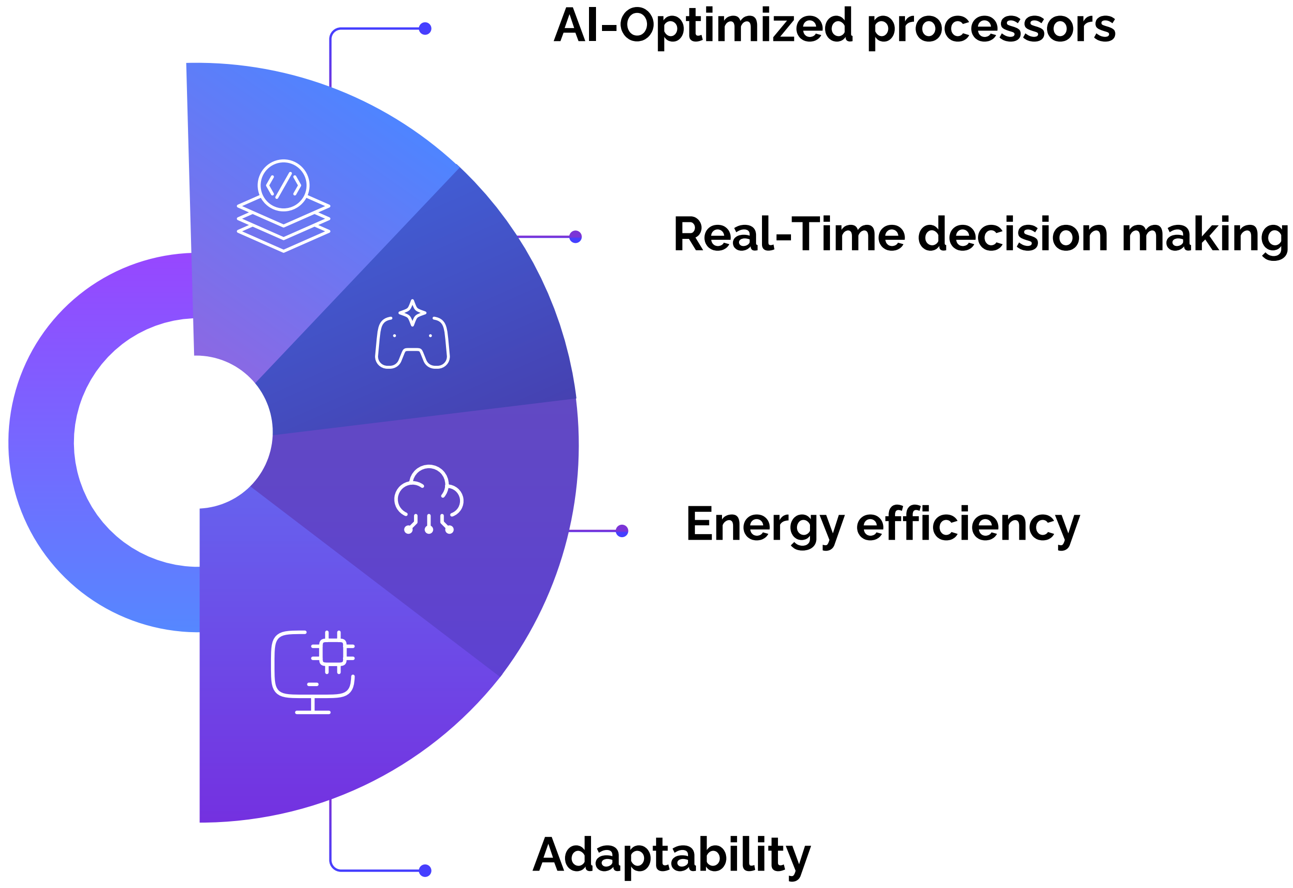

AI-Optimized processors:

- Dedicated AI cores (e.g., NPUs, TPUs) are built to handle neural network computations like matrix multiplication.
- Support for parallel processing enables faster deep learning inference and training.
- Chips are optimized for frameworks like TensorFlow, PyTorch, and ONNX, making them more AI-efficient.

Real-Time decision making:

- Smarter chips process data locally, reducing latency and making devices more responsive.
- On-device inference engines process data locally (e.g., Apple Neural Engine, NVIDIA Jetson).
- Eliminates round trips to the cloud, drastically reducing latency.
- Chips include real-time processing units (RTUs) that respond instantly to inputs (e.g., object detection in AR glasses).

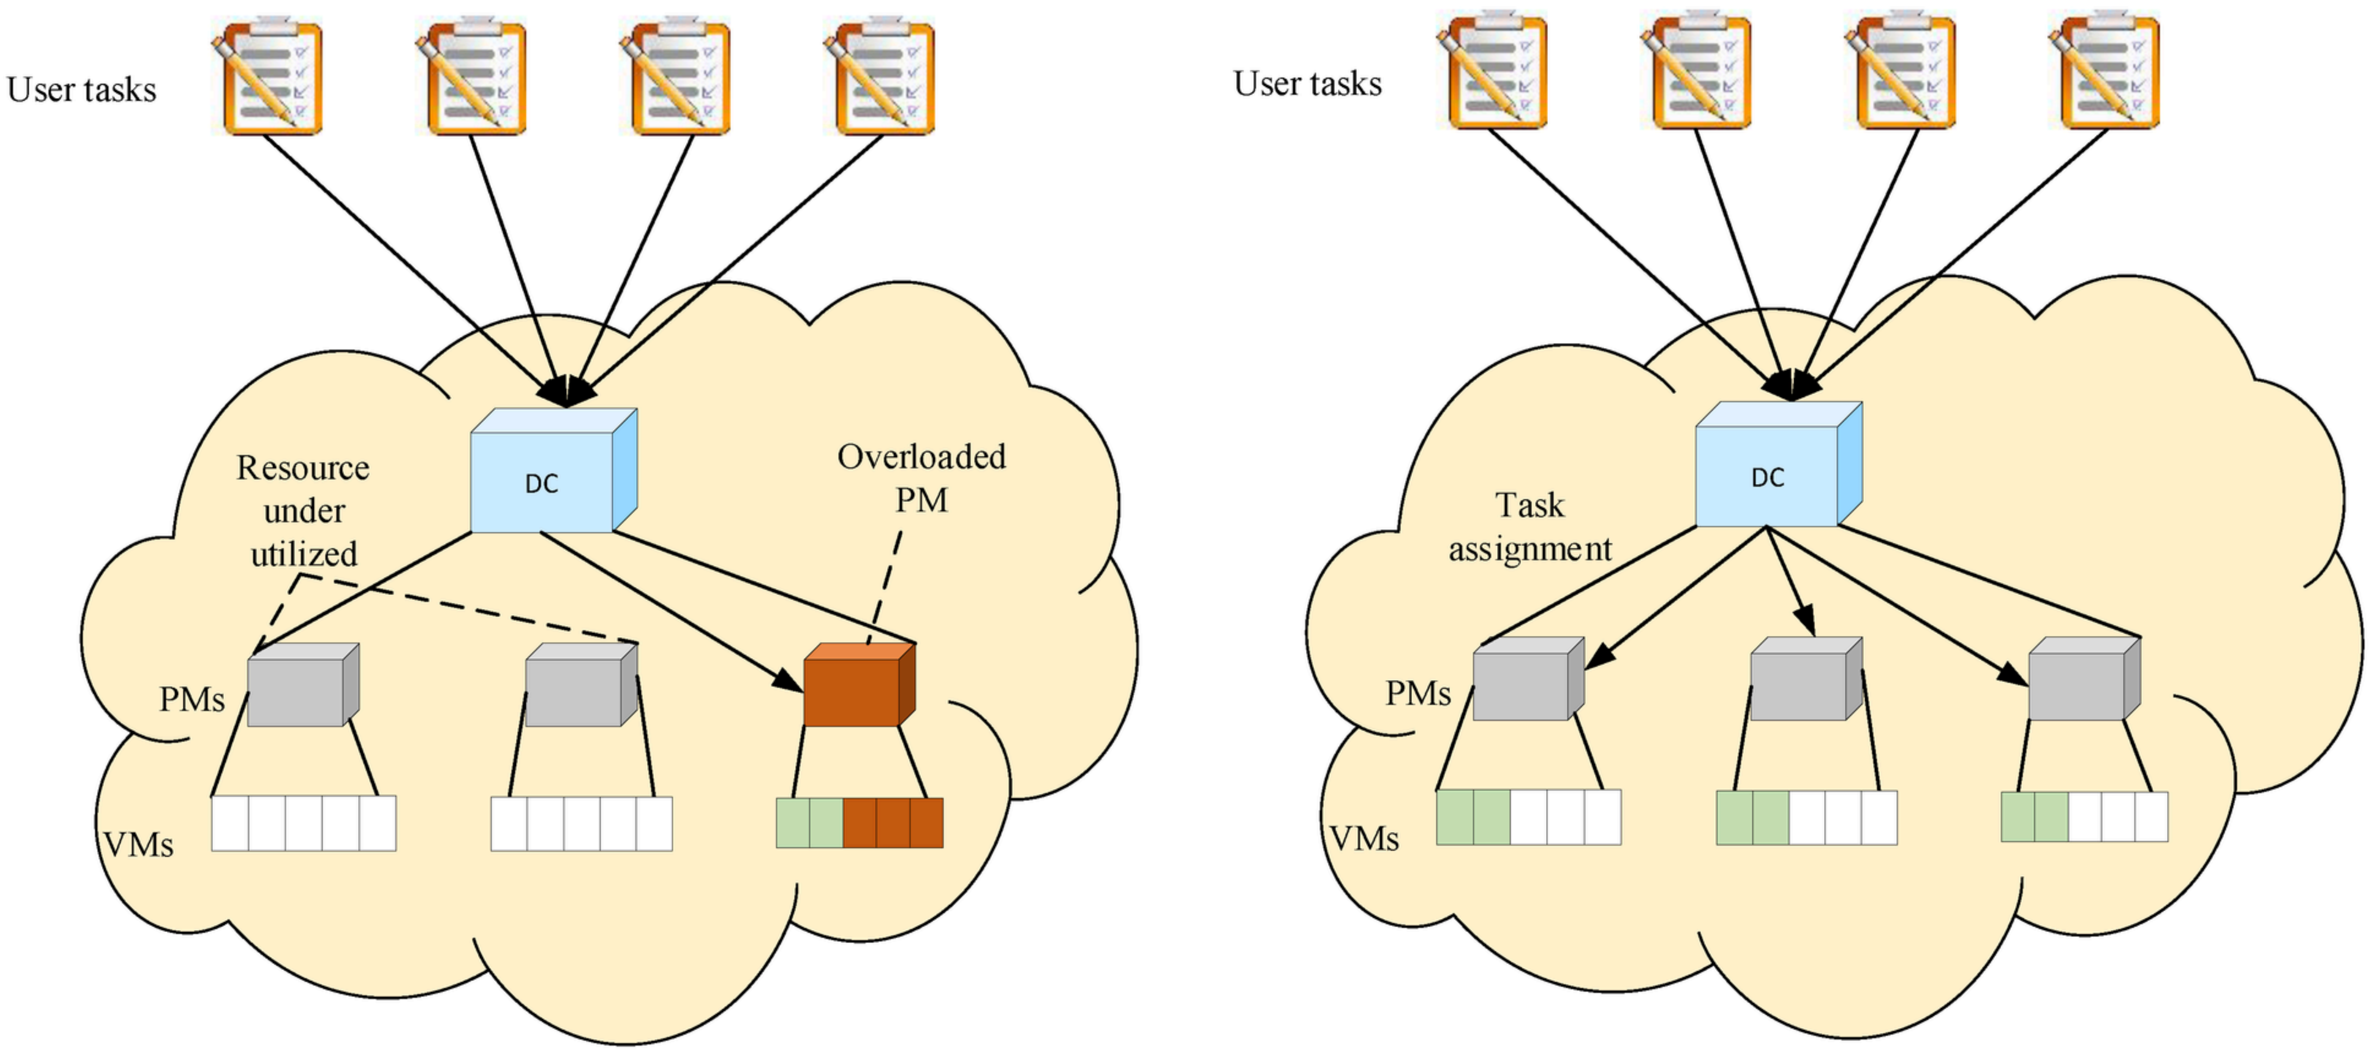

Block diagram for the overloaded and underloaded VMs.

Energy efficiency:

- Uses AI-based power scaling (e.g., adjusting clock speed and voltage based on task load).
- AI predicts usage patterns to optimize performance-per-watt.
- Some chips use low-power AI modes (like Qualcomm's AI Engine), which offload tasks from CPU to AI cores.

Adaptability:

- AI enables task-specific optimization e.g., allocating resources differently for voice vs. vision.
- Many processors now support modular AI models, dynamically loading and unloading tasks based on context.
- Chips like Intel Movidius and ARM Ethos-N adjust in real-time depending on workload complexity.

2. AI in Defense

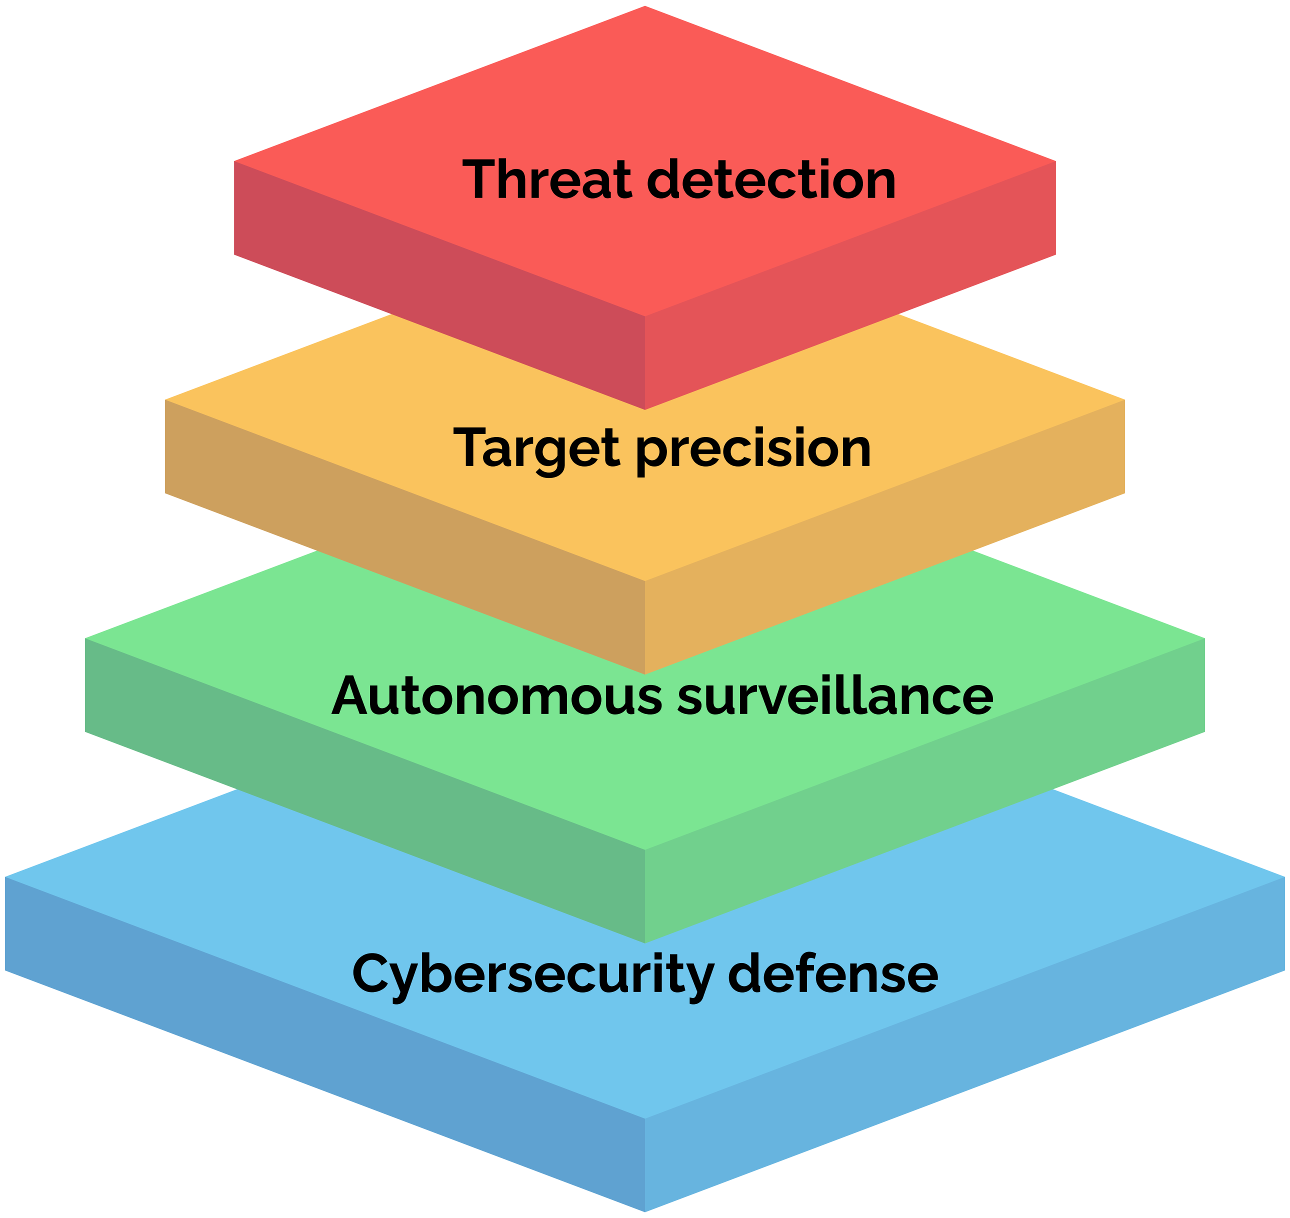

- **Threat detection:** CNNs/RNNs processes data from drones, satellites, and radars in real-time. Example: Project Maven's AI-hardware integration identifies missile launches in drone feeds, accelerating response.
- **Autonomous surveillance:** AI-driven drones (e.g., Israel's Harpy) pair YOLO algorithms with onboard AI processors to autonomously track and strike targets. Example: The U.S. Replicator Initiative deploys low-cost AI drones for swarm-based border monitoring.
- **Target precision:** Embedded AI chips in systems like F-35's Gorgon Strike use reinforcement learning to refine missile guidance, minimizing errors.
- **Cybersecurity defense:** AI-hardware hybrids like DARPA's "hacker bots" run Isolation Forest algorithms on specialized servers to block zero-day attacks in milliseconds.

3. AI + IoT Devices

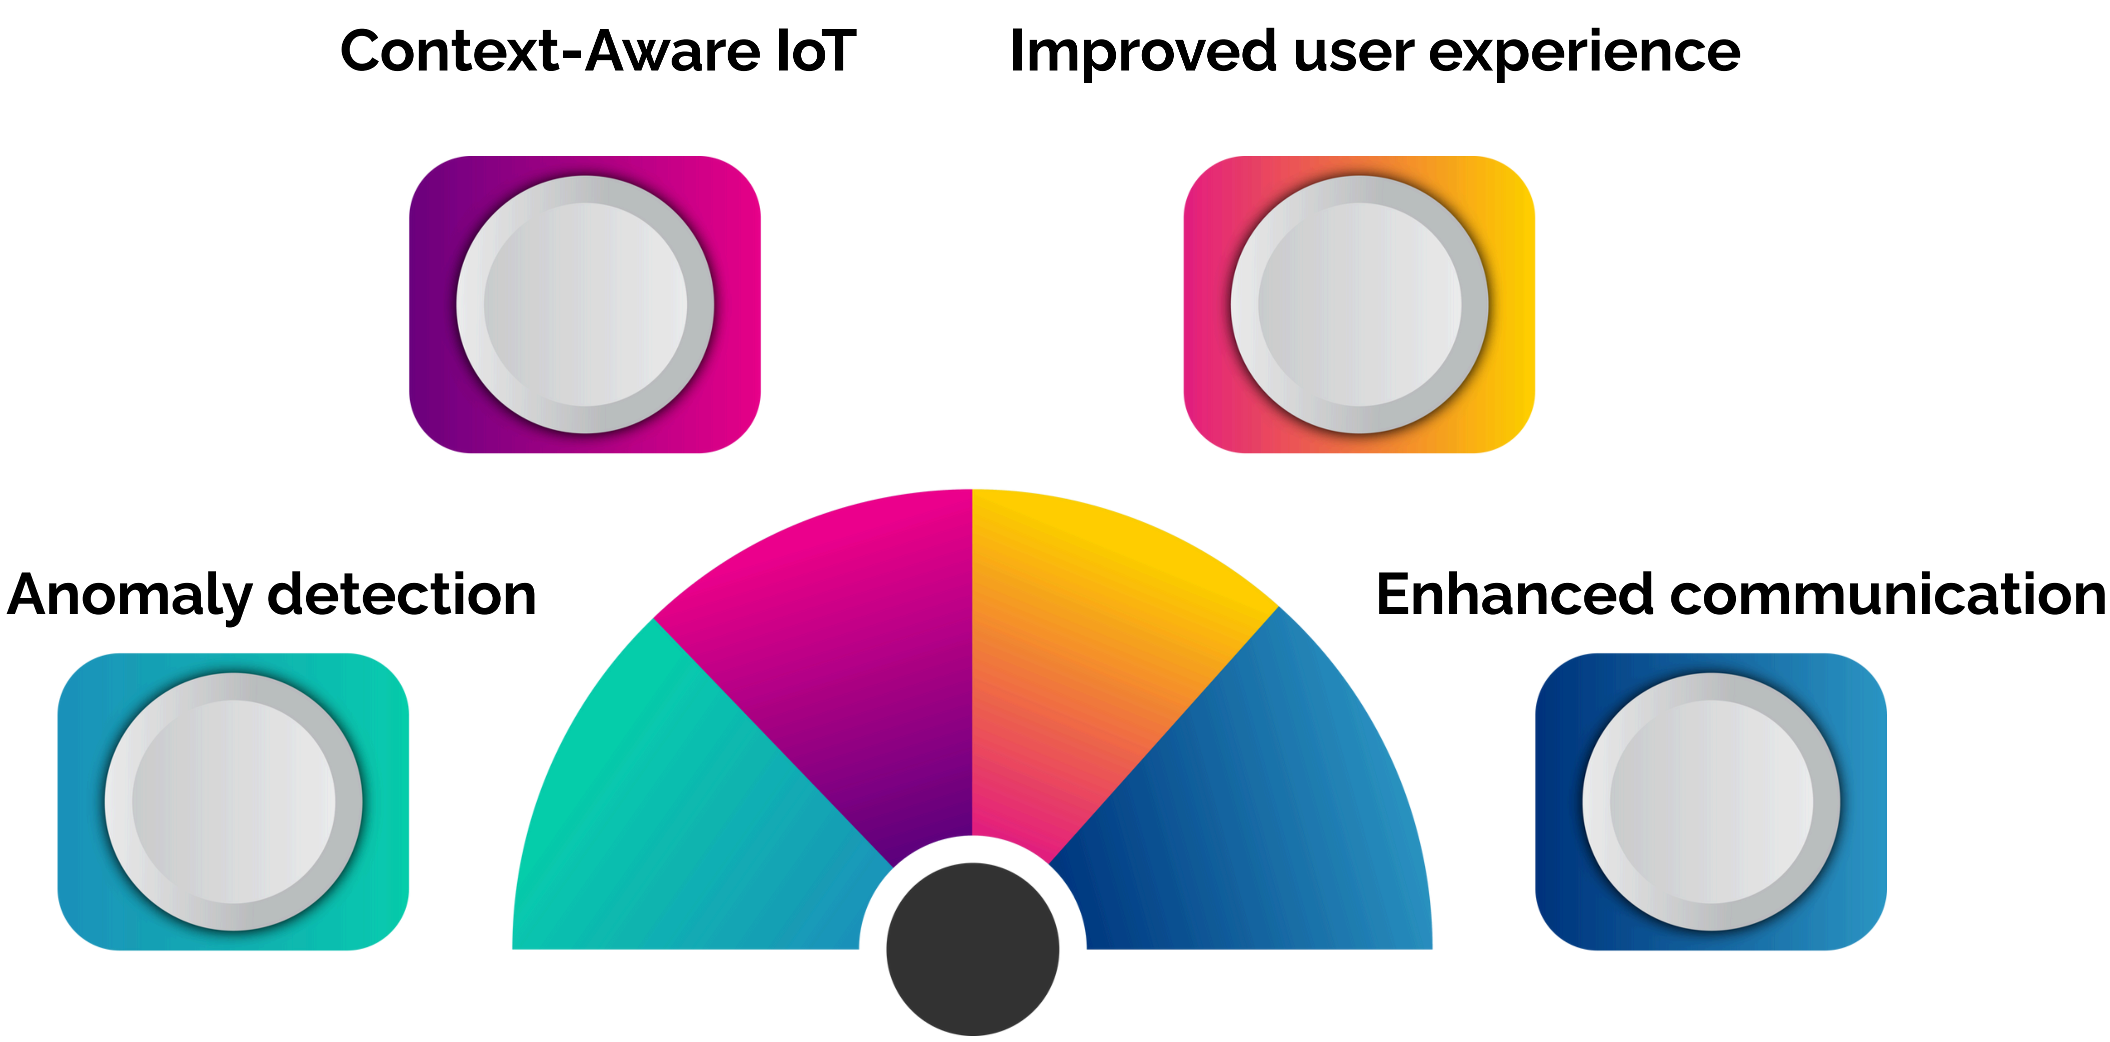

- **Context-aware IoT:** AI chips like Google's Edge TPU enable smart sensors to process environmental data locally. E.g: Nest Thermostats use RL on embedded ML hardware to adapt heating/cooling to user habits and occupancy.
- **Improved UX:** Amazon Echo's AZ1 Neural Engine run NLP models to personalize responses. Example: Smart lights (e.g., Philips Hue) learn routines via collaborative filtering algorithms on edge devices.
- **Anomaly detection:** AI-driven edge gateways (e.g., Siemens' IoT sensors) use autoencoders to flag factory equipment faults in real-time. Predictive maintenance in wind turbines pairs vibration sensors with on-device AI to prevent failures.
- **Enhanced communication:** NVIDIA's Jetson modules use federated learning to optimize IoT device coordination. Example: Smart cities like Barcelona deploy AI-hardware grids to sync traffic lights, reducing congestion by 20%.

4. AI + Edge Computing

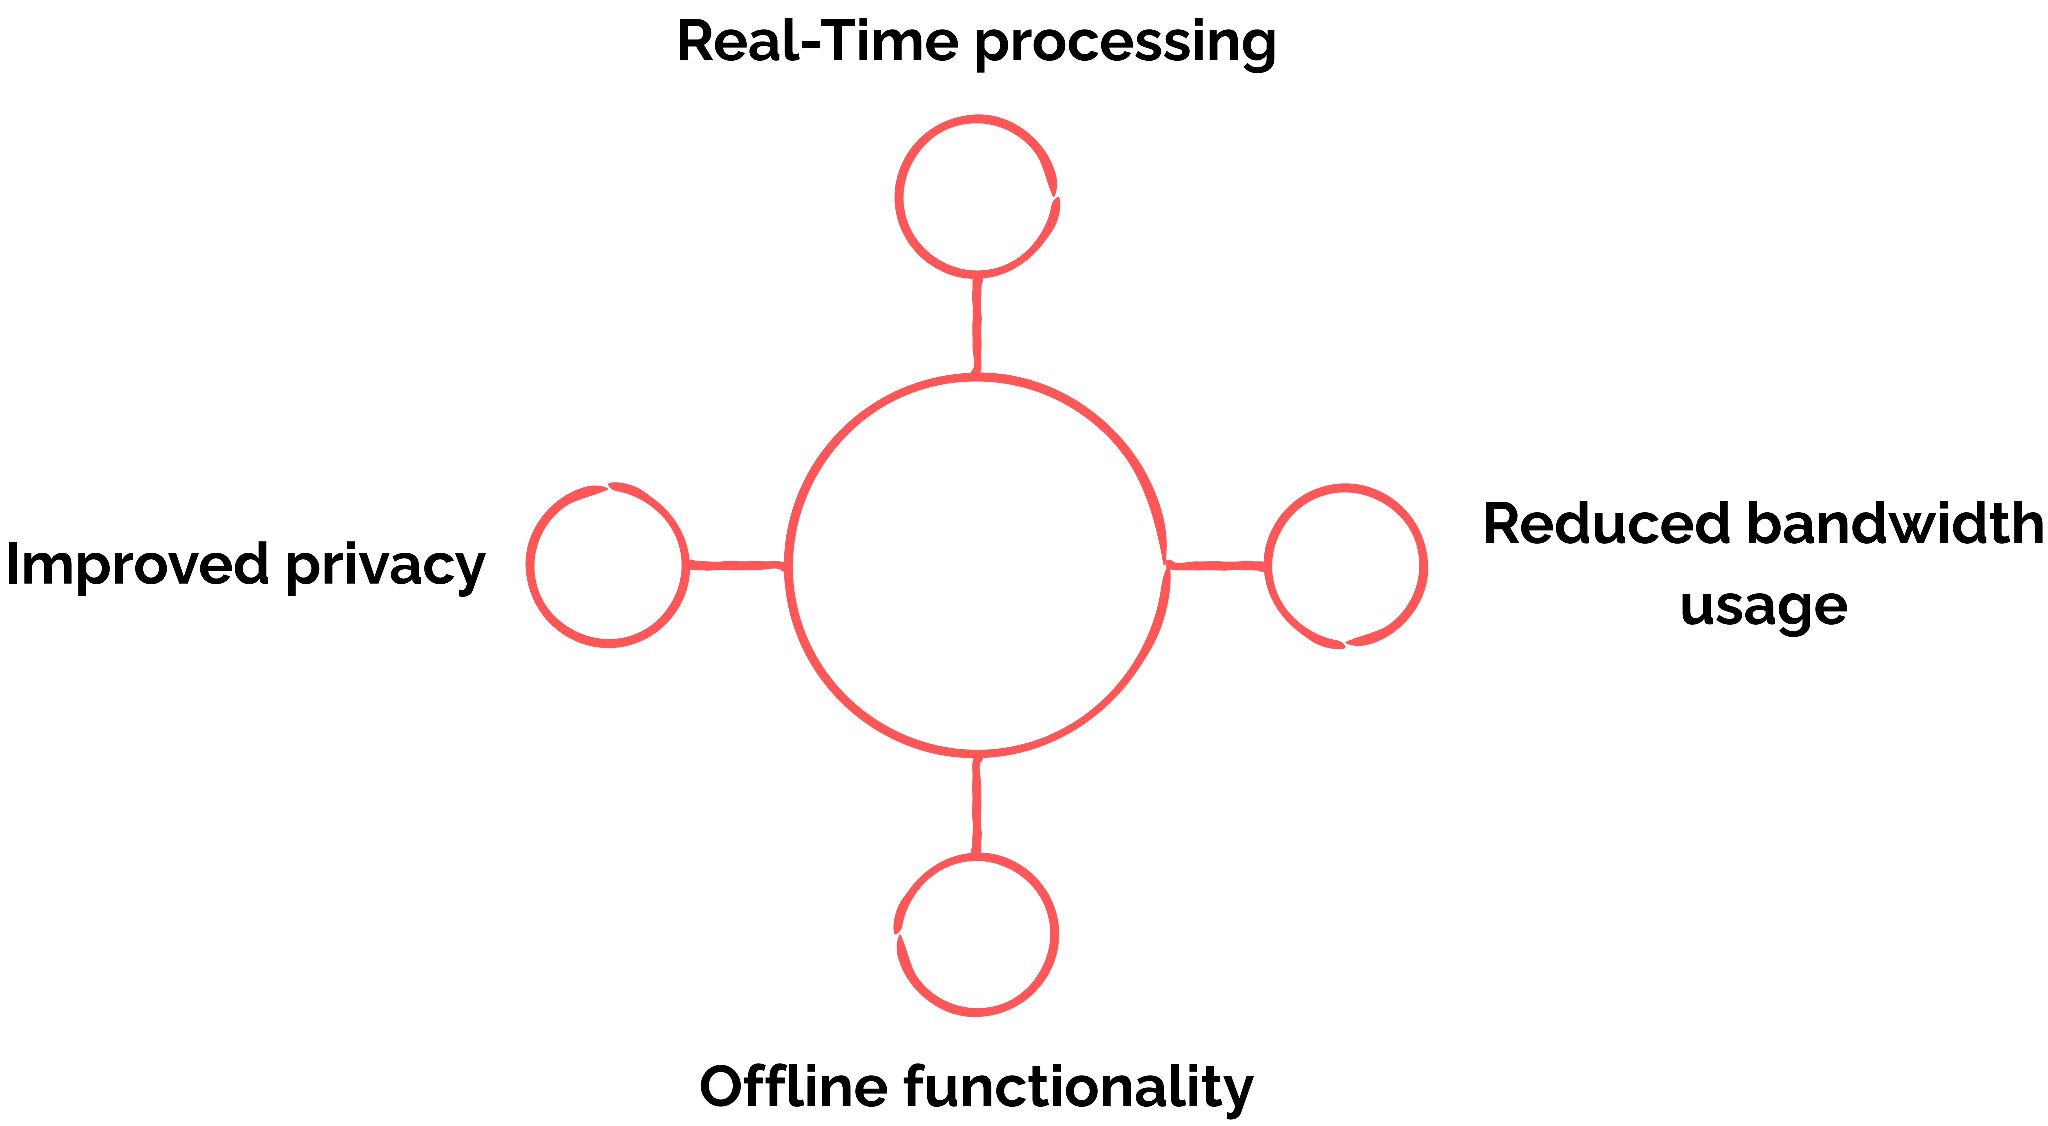

- **Real-Time processing:** Embedded AI chips run CNNs for instant decisions. Example: Autonomous vehicles process LiDAR data locally to avoid collisions.
- **Reduced bandwidth usage:** On-device AI (e.g., Intel Movidius VPUs) uses model pruning to filter non-critical data. E.g: Smart cameras send only flagged footage to the cloud.
- **Improved privacy:** Google Coral Edge TPUs enable federated learning for on-device data analysis. Example: Medical wearables process patient vitals locally, avoiding cloud exposure.
- **Offline functionality:** Qualcomm AI accelerators deploy reinforcement learning for autonomous drones in remote areas. Example: Agricultural drones analyze crop health offline via edge-based ML.

5. AI + Robotics

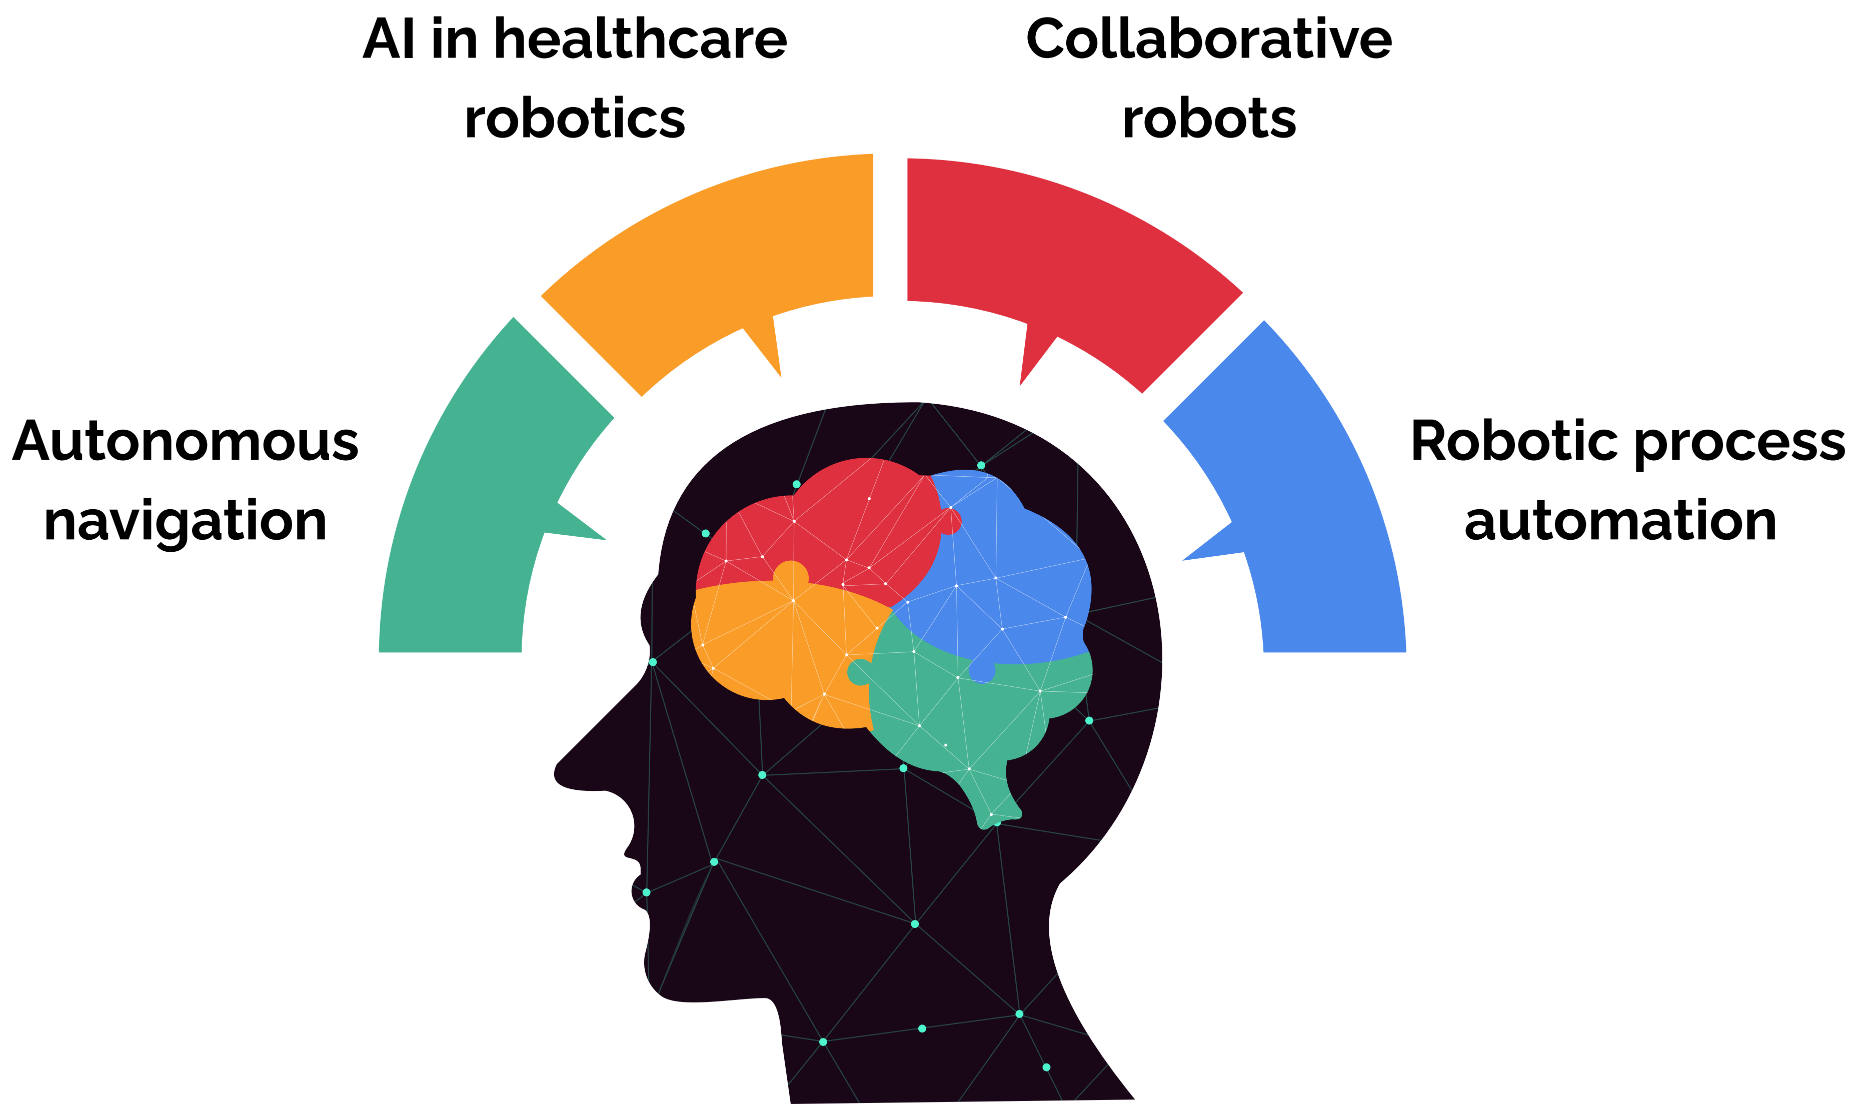

- **Autonomous navigation:** AI chips (e.g., NVIDIA Jetson) run CNNs and SLAM algorithms for obstacle avoidance. Example: Amazon's Proteus robots use real-time LiDAR + AI to navigate warehouses autonomously.
- **Robotic process automation (RPA):** Siemens PLCs integrate reinforcement learning to optimize assembly lines. Example: Foxconn's AI-driven arms use computer vision (YOLOv7) to sort electronics 3x faster.
- **Healthcare robotics:** Surgical bots like da Vinci leverage Kalman filters and RL for precision. Example: Intuitive Surgical systems reduce human error in prostate surgeries by 40% via AI-guided tools.
- **Collaborative robots :** Universal Robots' cobots pair force-sensitive algorithms with Azure AI Edge modules. Example: Toyota's cobots use federated learning to adapt to human workers' movements safely.

Stay Ahead with Our Tech Newsletter! 🚀

👉 **Subscribe now and never miss an update!**
🔗 https://bhavishyapandit9.substack.com/

Join our newsletter for:

- Step-by-step guides to mastering complex topics
- Industry trends & innovations delivered straight to your inbox
- Actionable tips to enhance your skills and stay competitive
- Insights on cutting-edge AI & software development

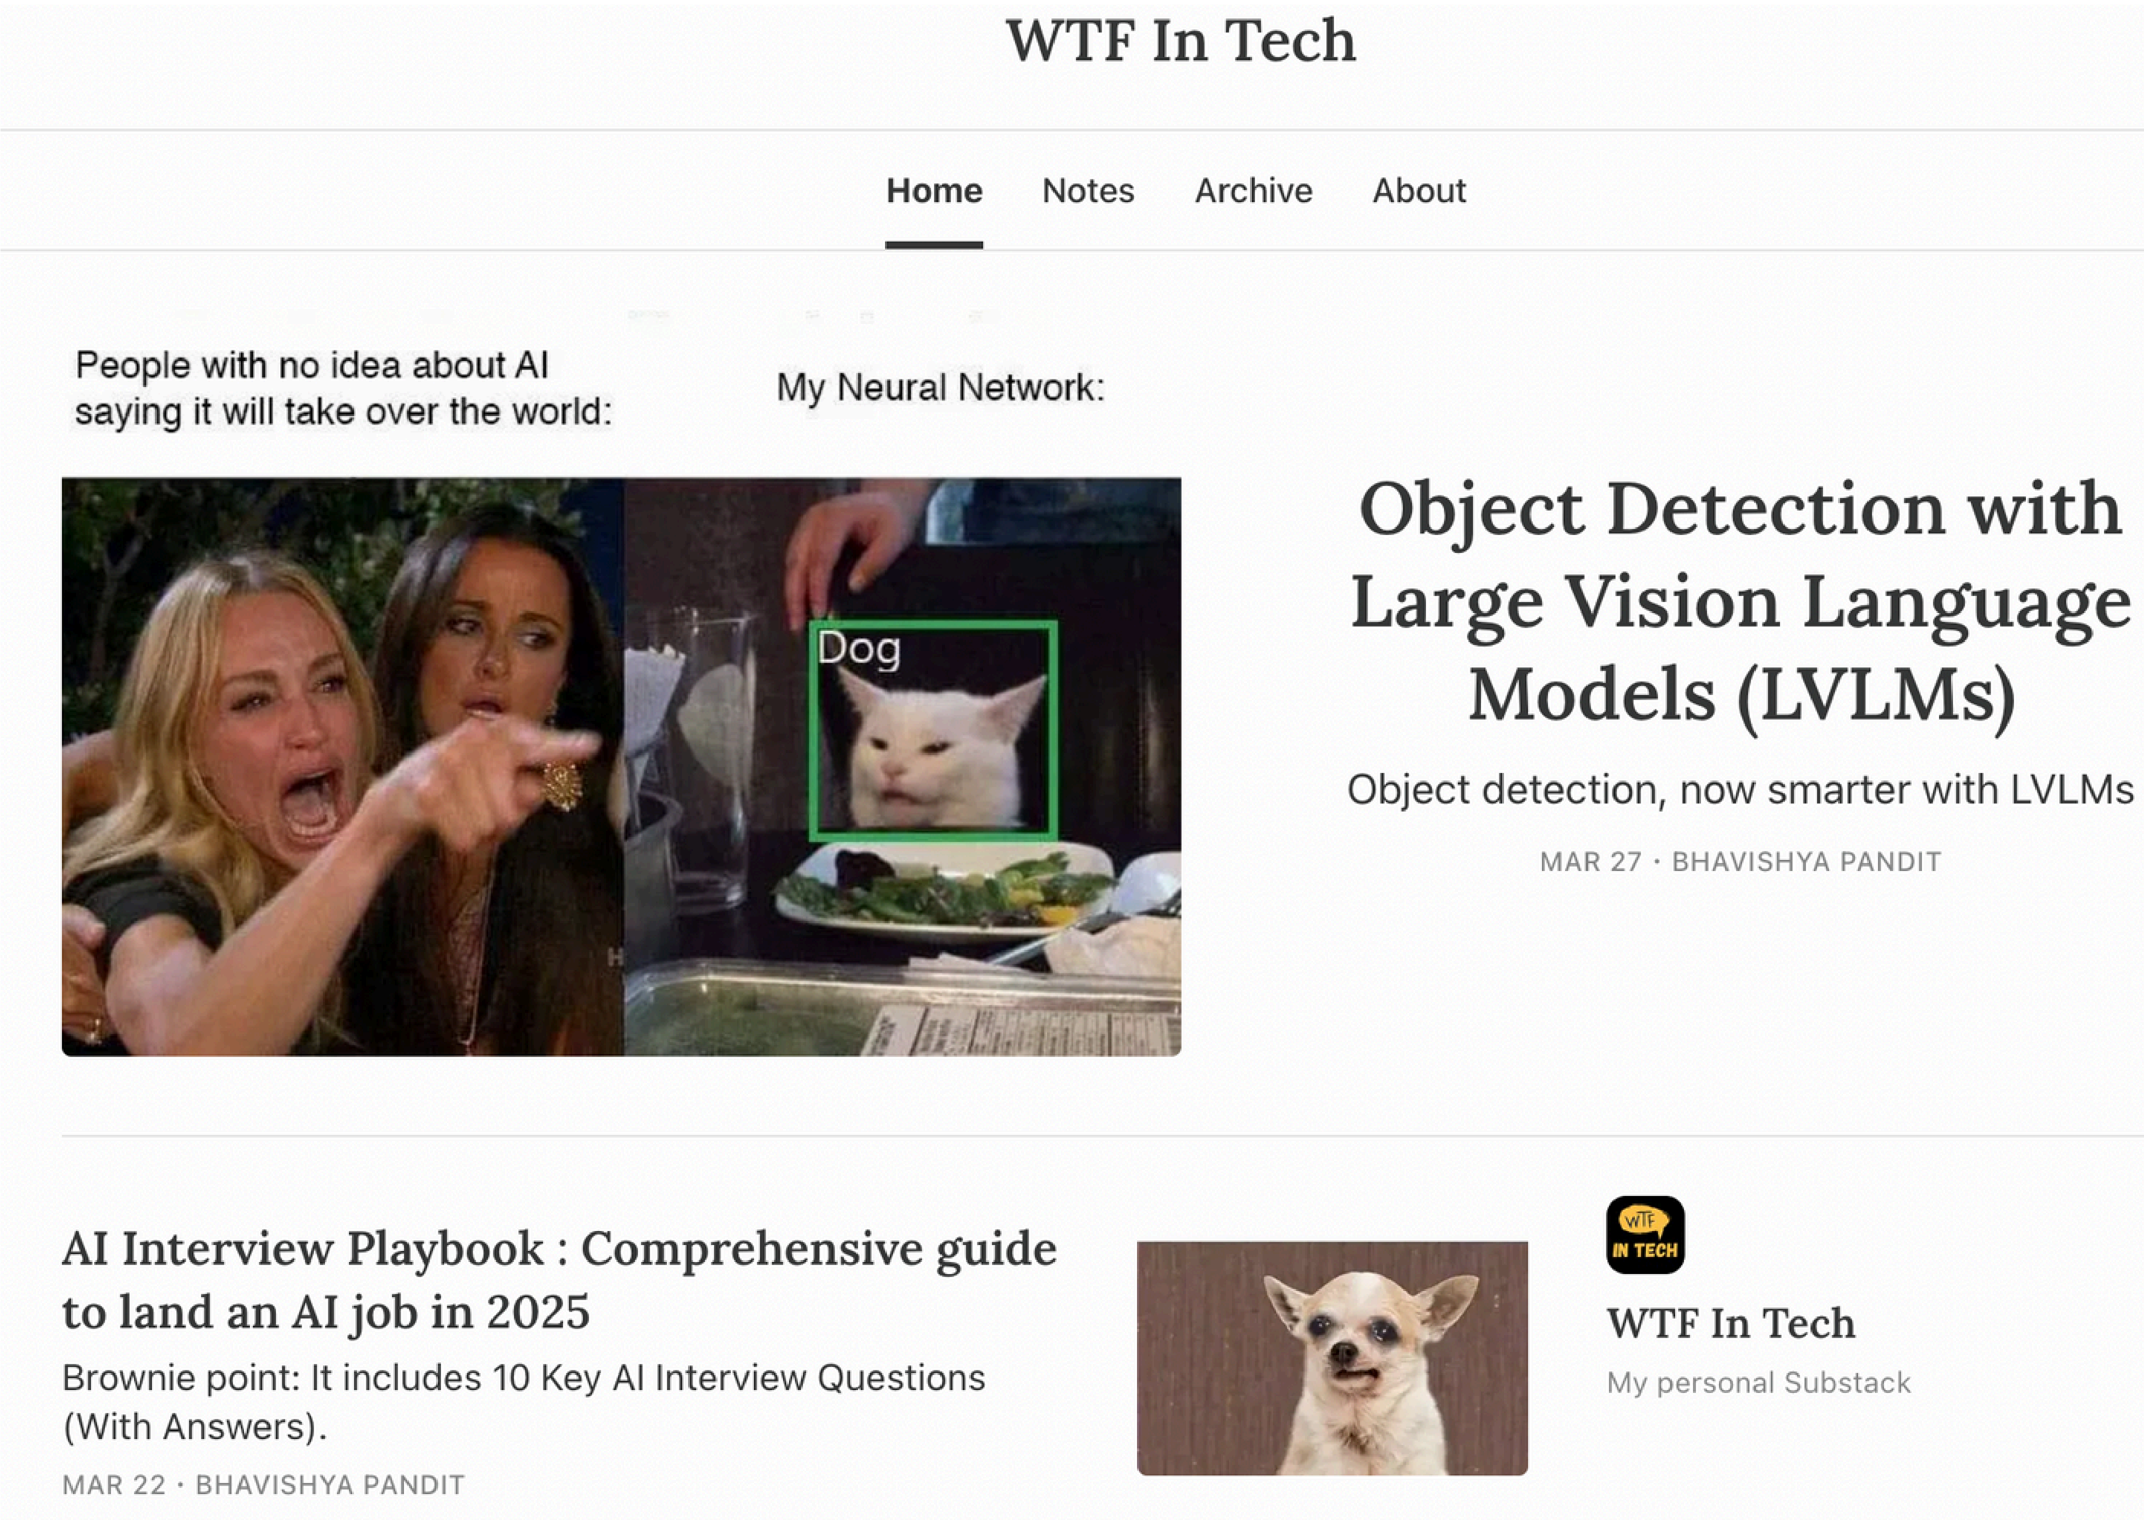

💡 Whether you're a developer, researcher, or tech enthusiast, this newsletter is your shortcut to staying informed and ahead of the curve.

Follow to stay updated on Generative AI

SAVE

LIKE

REPOST